

Biomedical Knowledge Graph Construction Using Large Language Models: Evaluating Generalisation Beyond Benchmark Datasets

Wesley Roots

September 2026

Contents

1	Introduction	2
2	Literature Review	3
2.0.1	Definition	3
2.0.2	Motivation	3
2.0.3	Tasks	3
2.0.4	Challenges	3
2.0.5	Traditional Methods	3
2.0.6	Transition to LLMs	3
3	Methodology	4
3.0.1	Research Design	4
3.0.2	Datasets	4
3.0.3	System Architecture	4
3.0.4	Information Extraction Pipeline	4
3.0.5	Pipeline Development	4
3.0.6	Contemporary Knowledge Graph Construction	4
3.0.7	Evaluation	4
3.0.8	Cost Analysis	4
3.0.9	Software Implementation	4
4	Results	5
5	Discussion	6
6	Conclusion	7

Chapter 1

Introduction

Chapter 2

Literature Review

2.0.1 Definition

Biomedical Information Extraction (IE) is the process of automatically identifying structured information - entities such as genes, diseases, chemicals and species, and the relationships between them - from unstructured biomedical text such as journal abstracts and full-text articles. This process is usually comprised of two main tasks: Named Entity Recognition (NER), which identifies biomedical entities, and Relation Extraction (RE), which identifies the relationships between those entities [1]. —

2.0.2 Motivation

The volume of biomedical literature has expanded at a rate that exceeds the capacity of manual review, with previous noncurrent analyses demonstrating an exponential growth in the number of biomedical publications [2]. This growth is supported by more recent data - a mapping of over 20 million PubMed abstracts has demonstrated the magnitude of contemporary biomedical literature to the level that the volume of biomedical and life-science publications has made it difficult for researchers to keep track of new literature [3]. This is not only a volume issue, but also an issue with discovery - log analysis of PubMed search behaviour shows that over 80% of views fall within the first page (top 20 results) of returned literature [4], meaning a substantial majority of literature goes largely unseen by researchers relying on manual review of ranked results.

2.0.3 Tasks

2.0.4 Challenges

2.0.5 Traditional Methods

2.0.6 Transition to LLMs

Chapter 3

Methodology

3.0.1 Research Design

3.0.2 Datasets

3.0.3 System Architecture

3.0.4 Information Extraction Pipeline

3.0.5 Pipeline Development

3.0.6 Contemporary Knowledge Graph Construction

3.0.7 Evaluation

3.0.8 Cost Analysis

3.0.9 Software Implementation

Chapter 4

Results

Chapter 5

Discussion

Chapter 6

Conclusion

Bibliography

- [1] Nadeesha Perera, Matthias Dehmer, and Frank Emmert-Streib. Named entity recognition and relation detection for biomedical information extraction. *Frontiers in cell and developmental biology*, 8:673, 2020.
- [2] Zhiyong Lu. Pubmed and beyond: a survey of web tools for searching biomedical literature. *Database*, 2011:baq036, 2011.
- [3] Rita González-Márquez, Luca Schmidt, Benjamin M Schmidt, Philipp Berens, and Dmitry Kobak. The landscape of biomedical research. *Patterns*, 5(6), 2024.
- [4] Rezarta Islamaj Dogan, G Craig Murray, Aurélie Névéol, and Zhiyong Lu. Understanding pubmed® user search behavior through log analysis. *Database*, 2009:bap018, 2009.