

Biomedical Knowledge Graph Construction Using Large Language Models: Evaluating Generalisation Beyond Benchmark Datasets

Wesley Roots

September 2026

Contents

1	Introduction	2
2	Literature Review	3
2.0.1	Definition	3
2.0.2	Motivation	3
2.0.3	Tasks	4
2.0.4	Challenges	4
2.0.5	Traditional Methods and Evaluation	4
2.0.6	Transition to LLMs	5
3	Methodology	6
3.0.1	Research Design	6
3.0.2	Datasets	6
3.0.3	System Architecture	6
3.0.4	Information Extraction Pipeline	6
3.0.5	Pipeline Development	6
3.0.6	Contemporary Knowledge Graph Construction	6
3.0.7	Evaluation	6
3.0.8	Cost Analysis	6
3.0.9	Software Implementation	6
4	Results	7
5	Discussion	8
6	Conclusion	9

Chapter 1

Introduction

Chapter 2

Literature Review

2.0.1 Definition

Biomedical Information Extraction (IE) is the process of automatically identifying structured information - entities such as genes, diseases, chemicals and species, and the relationships between them - from unstructured biomedical text such as journal abstracts and full-text articles. In this project, IE is treated as the first stage of a larger system: extracted entities and relations are later assembled into a knowledge graph and validated against biomedical ontologies [REFERENCE] (section)

2.0.2 Motivation

The volume of biomedical literature has expanded at a rate that exceeds the capacity of manual review, with previous noncurrent analyses demonstrating an exponential growth in the number of biomedical publications [1]. This growth is supported by more recent data - a mapping of over 20 million PubMed abstracts has demonstrated the magnitude of contemporary biomedical literature to the level that the volume of biomedical and life-science publications has made it difficult for researchers to keep track of new literature [2]. This is not only a volume issue, but also an issue with discovery - log analysis of PubMed search behaviour shows that over 80% of views fall within the first page (top 20 results) of returned literature [3], meaning a substantial majority of literature goes largely unseen by researchers relying on manual review of ranked results.

Together, these findings motivate the need for automated IE for two reasons: the breadth of literature is growing faster than it can be analysed [1, 2] and that there is a significant attention bottleneck that automated extraction and structuring could help mitigate [3]. This motivation supports the choice of PubMed as the source of the contemporary corpus for this thesis, as a corpus that is both large and continuously updated is exactly the setting in which static benchmark-trained pipelines are most likely to be used.

2.0.3 Tasks

Biomedical IE is usually comprised of several sub-tasks: Named Entity Recognition (NER), which identifies and stores types of spans of text referring to biomedical concepts, and Relation Extraction (RE), which identifies semantic relationships between recognised entities and Entity Normalisation (EN), which maps extracted mentions to pre-defined identifiers from an existing ontology [4, 5]. Long or complex sentences, which are common in biomedical abstracts, summarising multiple findings into a single sentence, can cause a significant challenge for these tasks [6].

2.0.4 Challenges

Beyond task complexity, biomedical text includes domain-specific challenges that separate it from general IE: dense and evolving terminology [7, 8], high rates of entity name ambiguity [9] (e.g.: synonymy such as "cerebral hemorrhage" and "intra-cerebral micro-bleeds" or, less-commonly, polysemy), and frequent use of abbreviations, many of which are ambiguous, defined within that paper, or polysemous.

In addition to these domain-specific challenges, significant methodological limitations compound the difficulty of relation extraction evaluation. Firstly, meaningful comparison between biomedical IE systems requires consistent evaluation frameworks, while strong generalisation across literature distributions requires models to perform reliably across heterogeneous data [REFERENCE].

2.0.5 Traditional Methods and Evaluation

Biomedical IE has progressed from rule-based systems towards feature-based machine learning and deep learning approaches built on pretrained transformer models. Benchmark-built systems have been at the core of this progression, with biomedical relation extraction models being commonly developed and evaluated using fixed, established benchmark corpora. . Recent exploration by Sanger et al. (2025) compares multiple pretrained language models before evaluating knowledge extraction of the best-performing model across five datasets with a consistent evaluation framework [10]. Both of these

Biomedical IE has progressed from rule-based systems towards feature-based machine learning and deep learning approaches built on pretrained transformer models. Throughout this progression, development and evaluation have remained largely benchmark-driven, with models commonly trained and evaluated on fixed, established corpora. However, reliance on individual benchmarks introduces the comparability and generalisability limitations outlined above. Recent work has therefore sought to address these limitations through more consistent evaluation frameworks and broader, heterogeneous training data. Sanger et al. (2025), for example, compare multiple pretrained language models before evaluating knowledge augmentation

of the best-performing model across five datasets within a consistent evaluation framework [10]. Similarly, Lai et al. (2023) introduce BioREx, which harmonises heterogeneous biomedical RE datasets to support broader training and improve generalisation [REFERENCE].

2.0.6 Transition to LLMs

Chapter 3

Methodology

3.0.1 Research Design

3.0.2 Datasets

3.0.3 System Architecture

3.0.4 Information Extraction Pipeline

3.0.5 Pipeline Development

3.0.6 Contemporary Knowledge Graph Construction

3.0.7 Evaluation

3.0.8 Cost Analysis

3.0.9 Software Implementation

Chapter 4

Results

Chapter 5

Discussion

Chapter 6

Conclusion

Bibliography

- [1] Zhiyong Lu. Pubmed and beyond: a survey of web tools for searching biomedical literature. *Database*, 2011:baq036, 2011.
- [2] Rita González-Márquez, Luca Schmidt, Benjamin M Schmidt, Philipp Berens, and Dmitry Kobak. The landscape of biomedical research. *Patterns*, 5(6), 2024.
- [3] Rezarta Islamaj Dogan, G Craig Murray, Aurélie Névéol, and Zhiyong Lu. Understanding pubmed® user search behavior through log analysis. *Database*, 2009:bap018, 2009.
- [4] Nadeesha Perera, Matthias Dehmer, and Frank Emmert-Streib. Named entity recognition and relation detection for biomedical information extraction. *Frontiers in cell and developmental biology*, 8:673, 2020.
- [5] Haodi Li, Qingcai Chen, Buzhou Tang, Xiaolong Wang, Hua Xu, Baohua Wang, and Dong Huang. Cnn-based ranking for biomedical entity normalization. *BMC bioinformatics*, 18(Suppl 11):385, 2017.
- [6] Isabel Segura-Bedmar, Paloma Martínez, and César de Pablo-Sánchez. A linguistic rule-based approach to extract drug-drug interactions from pharmacological documents. *BMC bioinformatics*, 12(Suppl 2):S1, 2011.
- [7] Michael Krauthammer and Goran Nenadic. Term identification in the biomedical literature. *Journal of biomedical informatics*, 37(6):512–526, 2004.
- [8] Marcos Da Silveira, Julio Cesar Dos Reis, and Cédric Pruski. Management of dynamic biomedical terminologies: current status and future challenges. *Yearbook of Medical informatics*, 24(01):125–133, 2015.
- [9] Sendong Zhao, Chang Su, Zhiyong Lu, and Fei Wang. Recent advances in biomedical literature mining. *Briefings in Bioinformatics*, 22(3):bbaa057, 2021.
- [10] Mario Sängler and Ulf Leser. Knowledge-augmented pre-trained language models for biomedical relation extraction. *BMC bioinformatics*, 26(1):240, 2025.